

Import Packages and Data

<https://www.kaggle.com/datasets/raminhuseyn/airline-customer-satisfaction>

```
In [89]: import os  
os.getcwd()
```

```
%cd "C:\Users\Angela\OneDrive\Desktop\ANA500"
```

```
C:\Users\Angela\OneDrive\Desktop\ANA500
```

```
In [137... #import packages  
import numpy as np  
import pandas as pd  
  
import seaborn as sns  
import matplotlib.pyplot as plt  
  
from sklearn.linear_model import LogisticRegression  
from sklearn.metrics import classification_report, confusion_matrix  
  
from sklearn.model_selection import train_test_split  
from sklearn.preprocessing import StandardScaler  
from sklearn.linear_model import LogisticRegression  
from sklearn import metrics  
from sklearn.metrics import classification_report
```

```
In [91]: #Load dataset  
df = pd.read_csv('Airline_customer_satisfaction.csv')
```

Prepare Data

```
In [92]: #Review data type of variables
```

```
df.dtypes
```

```
Out[92]: satisfaction      object
Customer Type      object
Age                int64
Type of Travel      object
Class              object
Flight Distance     int64
Seat comfort        int64
Departure/Arrival time convenient  int64
Food and drink      int64
Gate location       int64
Inflight wifi service  int64
Inflight entertainment  int64
Online support       int64
Ease of Online booking  int64
On-board service     int64
Leg room service     int64
Baggage handling     int64
Checkin service      int64
Cleanliness          int64
Online boarding      int64
Departure Delay in Minutes  int64
Arrival Delay in Minutes  float64
dtype: object
```

```
In [93]: #Review size of data
df.shape
```

```
Out[93]: (129880, 22)
```

```
In [94]: # check if there are null values
df.isnull().sum()
```

```
Out[94]: satisfaction      0
Customer Type            0
Age                     0
Type of Travel           0
Class                   0
Flight Distance          0
Seat comfort             0
Departure/Arrival time convenient  0
Food and drink           0
Gate location            0
Inflight wifi service    0
Inflight entertainment   0
Online support            0
Ease of Online booking   0
On-board service         0
Leg room service         0
Baggage handling         0
Checkin service          0
Cleanliness              0
Online boarding           0
Departure Delay in Minutes  0
Arrival Delay in Minutes  393
dtype: int64
```

Arrival Delay in Minutes has 393 missing values.

Handle Missing Values

```
In [66]: #Take a closer Look at Arrival Delay in Minutes
df['Arrival Delay in Minutes'].describe()
```

```
Out[66]: count      129487.000000
         mean        15.091129
         std         38.465650
         min          0.000000
         25%          0.000000
         50%          0.000000
         75%         13.000000
         max        1584.000000
         Name: Arrival Delay in Minutes, dtype: float64
```

Arrival Delay in Minutes has values that range from 0 to 1584. The values have a mean of 15.091129

```
In [95]: #Create new variable with copy of original data
Arrival_Delay = df['Arrival Delay in Minutes']

#Add new variable to dataframe
df2 = df.assign(Arrival_Delay_Minutes = Arrival_Delay)

#Check that new variable matches original variable
df2['Arrival_Delay_Minutes'].describe()
```

```
Out[95]: count      129487.000000
         mean        15.091129
         std         38.465650
         min          0.000000
         25%          0.000000
         50%          0.000000
         75%         13.000000
         max        1584.000000
         Name: Arrival_Delay_Minutes, dtype: float64
```

I want to ensure that I keep the integrity of the original data. Therefore, I created a new variable to use in place of Arrival Delay in Minutes. I will be copying Arrival Delay in Minutes to the new variable. I will be manipulating the new variable to handle the missing values.

```
In [68]: #fill missing values on new variable with mean value
# filling missing value using fillna()
df2['Arrival_Delay_Minutes'].fillna(15.091129, inplace = True)
```

```
#check the new variable to see if the missing values are updated
df2['Arrival_Delay_Minutes'].isnull().sum()
```

Out[68]: 0

Because there are a small amount of missing values, I chose to replace them with the Mean value of the variable. This will help keep more data to work with rather than dropping the values.

```
In [96]: #Drop the original variable in second dataframe
df2.drop('Arrival Delay in Minutes', axis=1, inplace=True)

#Check that the variable was dropped
df2.columns
```

Out[96]: Index(['satisfaction', 'Customer Type', 'Age', 'Type of Travel', 'Class',
 'Flight Distance', 'Seat comfort', 'Departure/Arrival time convenient',
 'Food and drink', 'Gate location', 'Inflight wifi service',
 'Inflight entertainment', 'Online support', 'Ease of Online booking',
 'On-board service', 'Leg room service', 'Baggage handling',
 'Checkin service', 'Cleanliness', 'Online boarding',
 'Departure Delay in Minutes', 'Arrival_Delay_Minutes'],
 dtype='object')

Since the new variable was assigned to a new dataframe, I chose to drop Arrival Delay in Minutes since the data is duplicative of the new variable Arrival_Delay_Minutes.

Data Visualizations

```
In [97]: #Print the statistics of the variables
df2.describe()
```

Out[97]:

	Age	Flight Distance	Seat comfort	Departure/Arrival time convenient	Food and drink	Gate location	Inflight wifi service	e
count	129880.000000	129880.000000	129880.000000	129880.000000	129880.000000	129880.000000	129880.000000	1
mean	39.427957	1981.409055	2.838597	2.990645	2.851994	2.990422	3.249130	
std	15.119360	1027.115606	1.392983	1.527224	1.443729	1.305970	1.318818	
min	7.000000	50.000000	0.000000	0.000000	0.000000	0.000000	0.000000	
25%	27.000000	1359.000000	2.000000	2.000000	2.000000	2.000000	2.000000	
50%	40.000000	1925.000000	3.000000	3.000000	3.000000	3.000000	3.000000	
75%	51.000000	2544.000000	4.000000	4.000000	4.000000	4.000000	4.000000	
max	85.000000	6951.000000	5.000000	5.000000	5.000000	5.000000	5.000000	

Numeric Variables


```
In [98]: #Create dataframe with only numeric variables
num_df = df2.select_dtypes(include = np.number)
num_df.columns = num_df.columns.str.replace(' ', '_')
num_df.columns = num_df.columns.str.replace('/', '_')
num_df.columns = num_df.columns.str.replace('-', '_')

#Check to see that the dataframe is correct
num_df.dtypes
```

```
Out[98]: Age int64
Flight_Distance int64
Seat_comfort int64
Departure_Arrival_time_convenient int64
Food_and_drink int64
Gate_location int64
Inflight_wifi_service int64
Inflight_entertainment int64
Online_support int64
Ease_of_Online_booking int64
On_board_service int64
Leg_room_service int64
Baggage_handling int64
Checkin_service int64
Cleanliness int64
Online_boarding int64
Departure_Delay_in_Minutes int64
Arrival_Delay_Minutes float64
dtype: object
```


```
In [99]: #Age Histogram
sns.histplot(num_df.Age, bins=30, kde=True, color='grey', edgecolor='purple')
plt.title('Age')
```

```
Out[99]: Text(0.5, 1.0, 'Age')
```



```
In [100... sns.boxplot(x= num_df.Age)
```


```
Out[100]: <Axes: xlabel='Age'>
```

Customers age ranges from 10 years old to 80+ years old. Age has a normal distribution. Age does not have any outliers. The most ages that travel are between 20 and 50 years old.

```
In [101... #Flight Distance Histogram  
sns.histplot(num_df.Flight_Distance, bins=30, kde=True, color='grey', edgecolor='purple')  
plt.title('Flight Distance')
```


```
Out[101]: Text(0.5, 1.0, 'Flight Distance')
```



There appears to be a right skew for Flight Distance. The data appears to go flat at 6000. It also appears that there may be some outlier at 7000.

```
In [102... #Flight Distance BoxPlot  
sns.boxplot(x= num_df.Flight_Distance)
```

```
Out[102]: <Axes: xlabel='Flight_Distance'>
```



Outliers are shown just after 4000.

In [103...

```
# How many values are outliers?  
  
flight_outliers = num_df[(num_df['Flight_Distance'] > 4200)]  
flight_outliers.Flight_Distance.describe()
```

```
Out[103]: count      3011.000000
          mean      4813.997675
          std       482.818474
          min       4201.000000
          25%       4429.500000
          50%       4729.000000
          75%       5103.000000
          max       6951.000000
          Name: Flight_Distance, dtype: float64
```

There are only 3011 values that are greater than 4200 which are listed as outliers on the boxplot. These values will be dropped since they are such a small portion of the data.

In [104...

```
def removal_box_plot(df, column, threshold):
    sns.boxplot(df[column])
    plt.title(f'Original Box Plot of {column}')
    plt.show()


    removed_outliers = df[df[column] <= threshold]

    sns.boxplot(removed_outliers[column])
    plt.title(f'Box Plot without Outliers of {column}')
    plt.show()
    return removed_outliers


threshold_value = 4200

no_outliers = removal_box_plot(num_df, 'Flight_Distance', threshold_value)
```

Original Box Plot of Flight_Distance



Box Plot without Outliers of Flight_Distance



In [105...


```
# drop rows containing outliers  
df3 = df2.drop(flight_outliers.index)  
df3['Flight Distance'].describe()
```

```
Out[105]: count    126869.000000
          mean      1914.182826
          std        937.831785
          min         50.000000
          25%       1337.000000
          50%       1900.000000
          75%       2488.000000
          max       4200.000000
          Name: Flight Distance, dtype: float64
```

After Flight Distance outliers are removed, there are 126,869 values remaining. The values range from 50 to 4,200.

```
In [106... # Histogram of Seat Comfort
sns.histplot(num_df.Seat_comfort, bins=30, color='grey', edgecolor='purple')
plt.title('Seat Comfort')
```

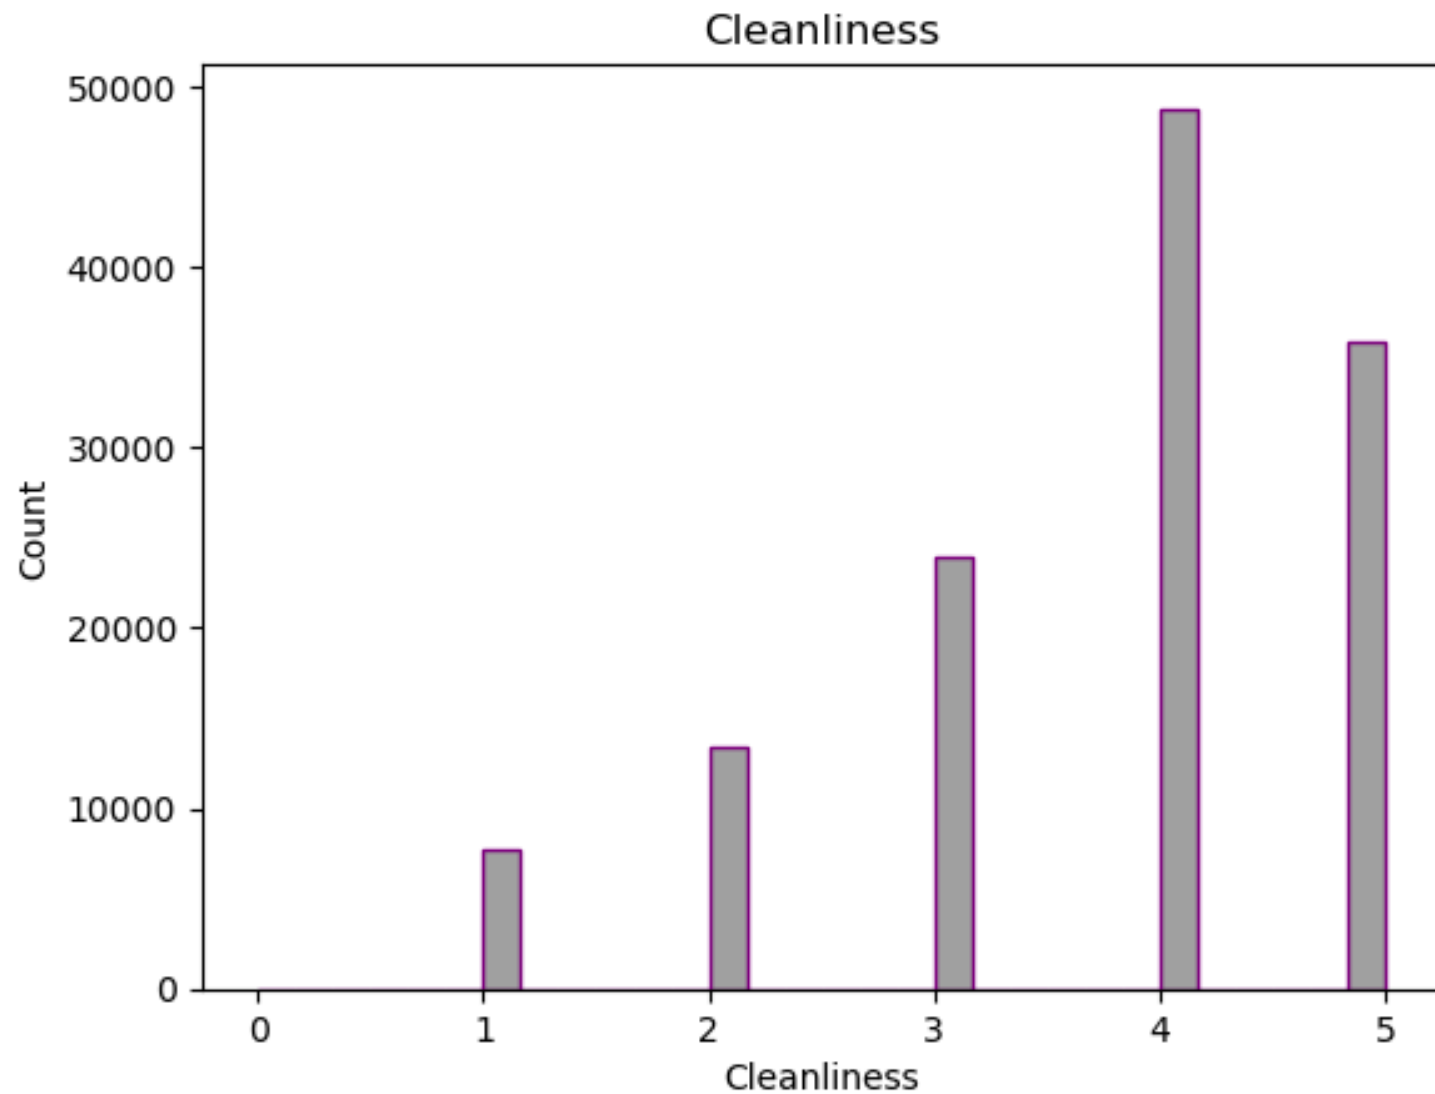
```
Out[106]: Text(0.5, 1.0, 'Seat Comfort')
```



Customers age ranges from 10 years old to 80+ years old. Age has a normal distribution. Age does not have any outliers. The most ages that travel are between 20 and 50 years old.

```
In [211... #Cleanliness Histogram
sns.histplot(num_df.Cleanliness, bins=30, color='grey', edgecolor='purple')
plt.title('Cleanliness')
```

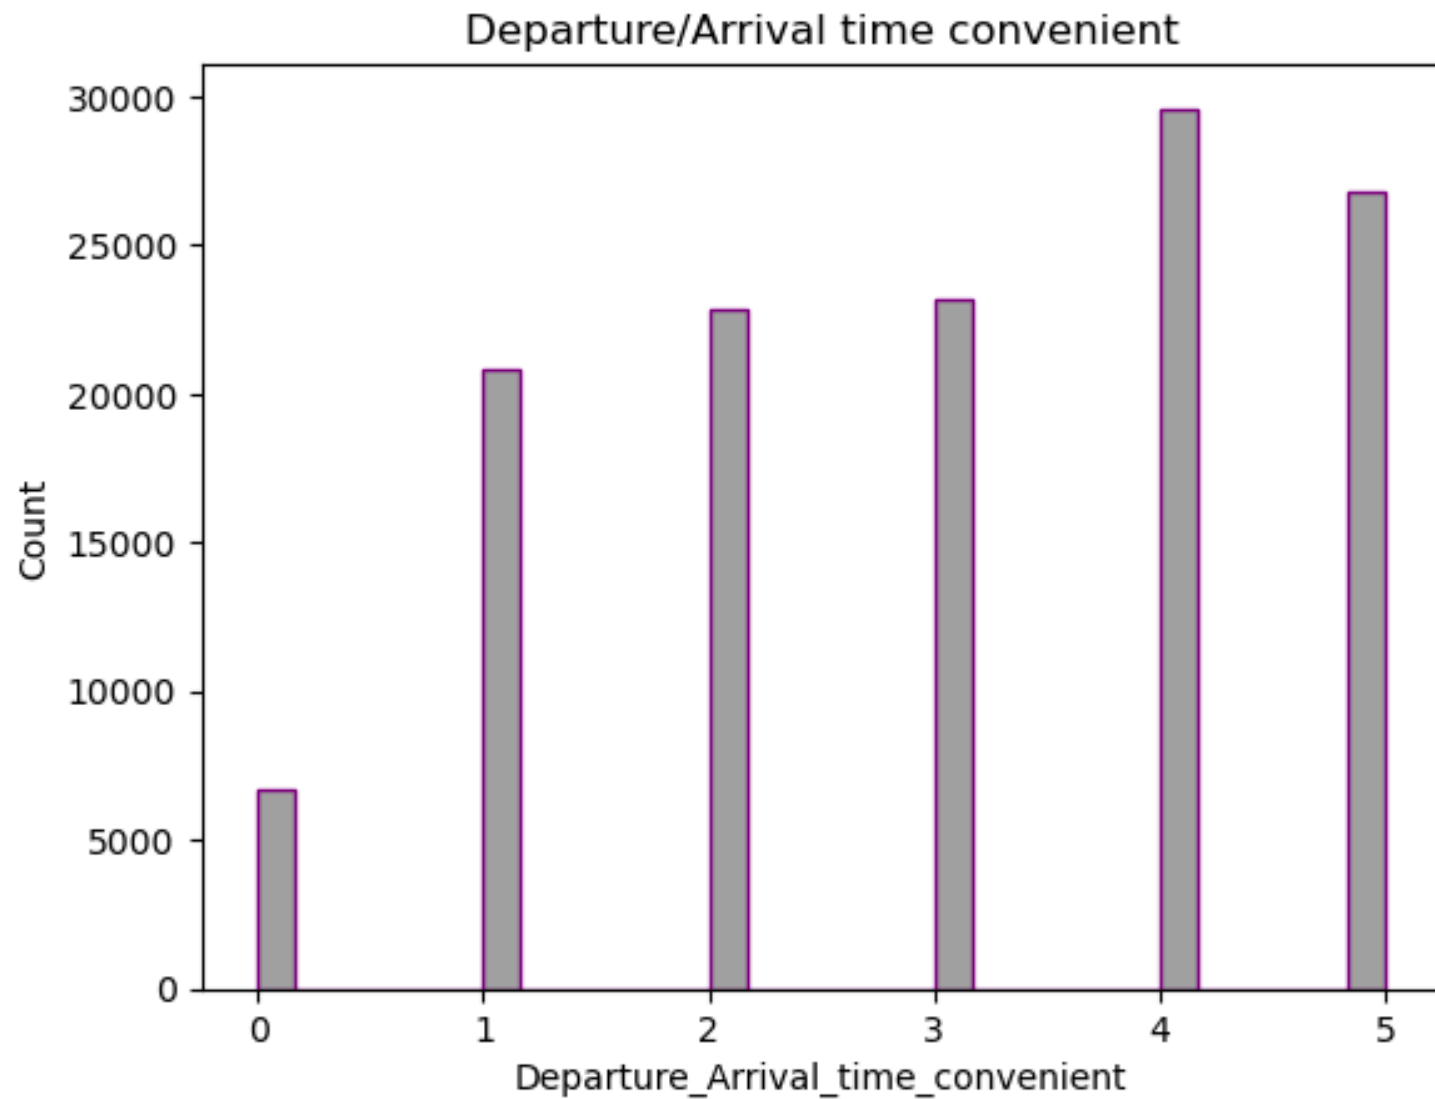
```
Out[211]: Text(0.5, 1.0, 'Cleanliness')
```

Cleanliness is a categorical variable that ranges from 0 to 5. The rating of 4 has the highest count of customers.

```
In [212... #Departure/Arrival time convenient Histogram
sns.histplot(num_df.Departure_Arrival_time_convenient, bins=30, color='grey', edgecolor='purple')
plt.title('Departure/Arrival time convenient')
```

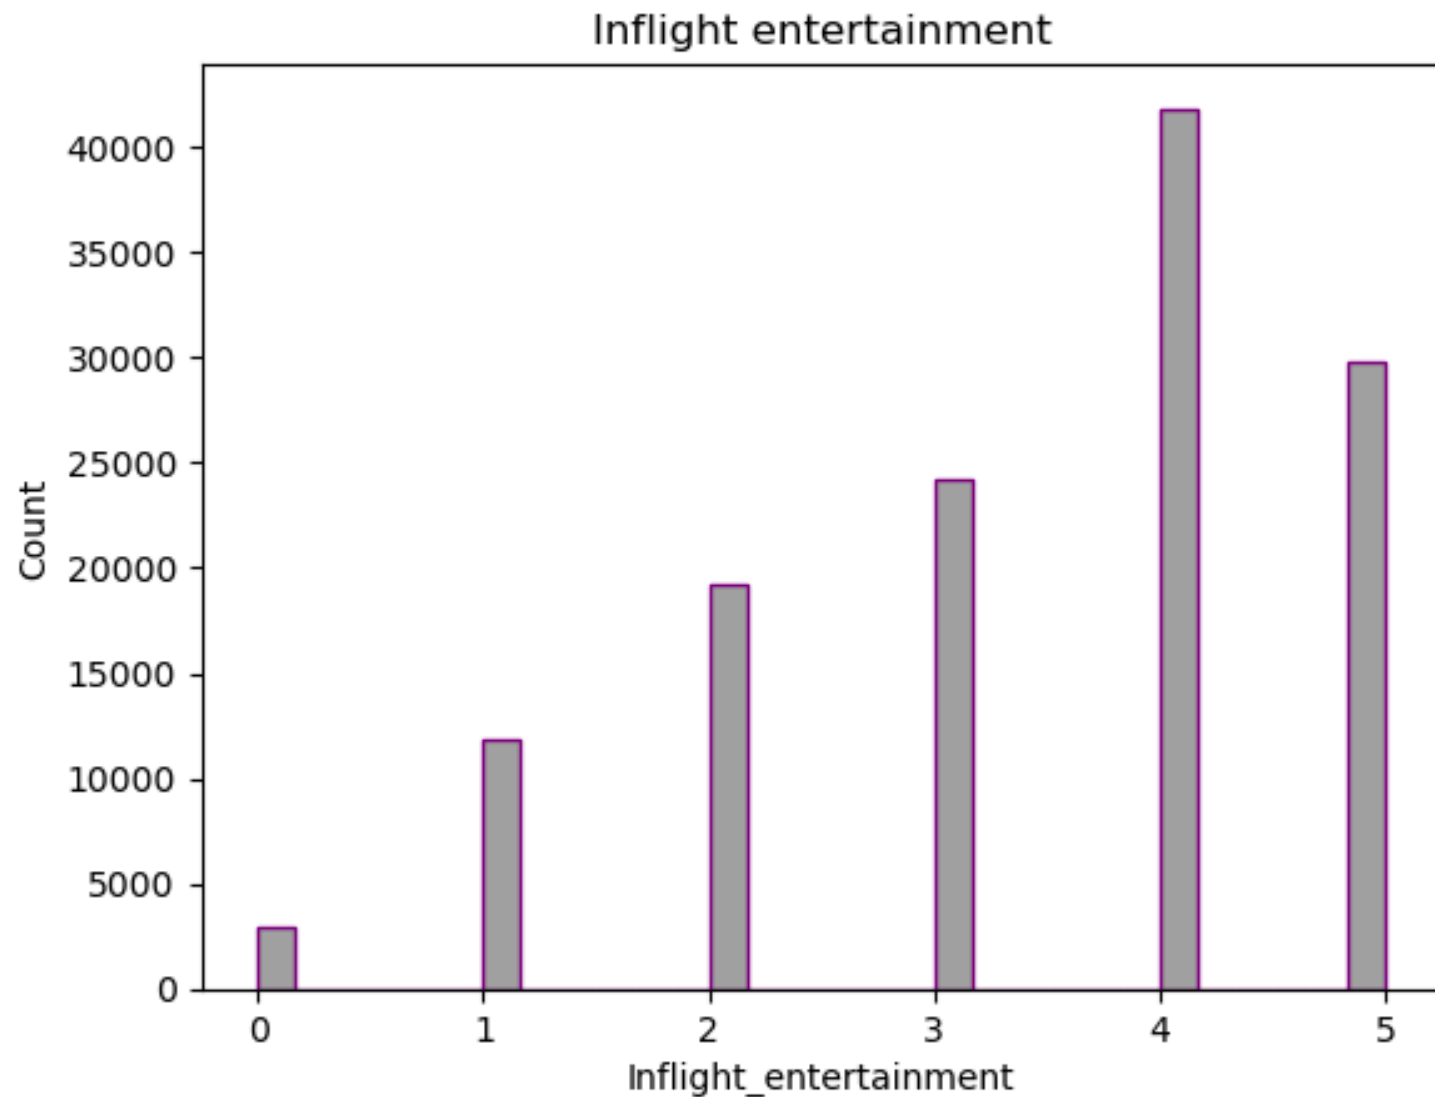
```
Out[212]: Text(0.5, 1.0, 'Departure/Arrival time convenient')
```



Departure/Arrival time convenient is a categorical variable. The rating scale ranges from 0 to 5. The rating 4 has the highest count of customers.

```
In [213... #Inflight entertainment Histogram
sns.histplot(num_df.Inflight_entertainment, bins=30, color='grey', edgecolor='purple')
plt.title('Inflight entertainment')
```

```
Out[213]: Text(0.5, 1.0, 'Inflight entertainment')
```



Inflight entertainment is a categorical variable. Inflight Entertainment is a rating scale from 0 to 5. The rating 4 has the highest count of customers.

```
In [215... #Inflight WiFi service Histogram  
sns.histplot(num_df.Inflight_wifi_service, bins=30, color='grey', edgecolor='purple')  
plt.title('Inflight WiFi service')
```


```
Out[215]: Text(0.5, 1.0, 'Inflight WiFi service')
```



Inflight WiFi service is a categorical variable. Inflight WiFi service is a rating scale from 0 to 5. The rating 4 has the highest count of customers.

```
In [217... #Food and Drink Histogram  
sns.histplot(num_df.Food_and_drink, bins=30, color='grey', edgecolor='purple')  
plt.title('Food and Drink')
```

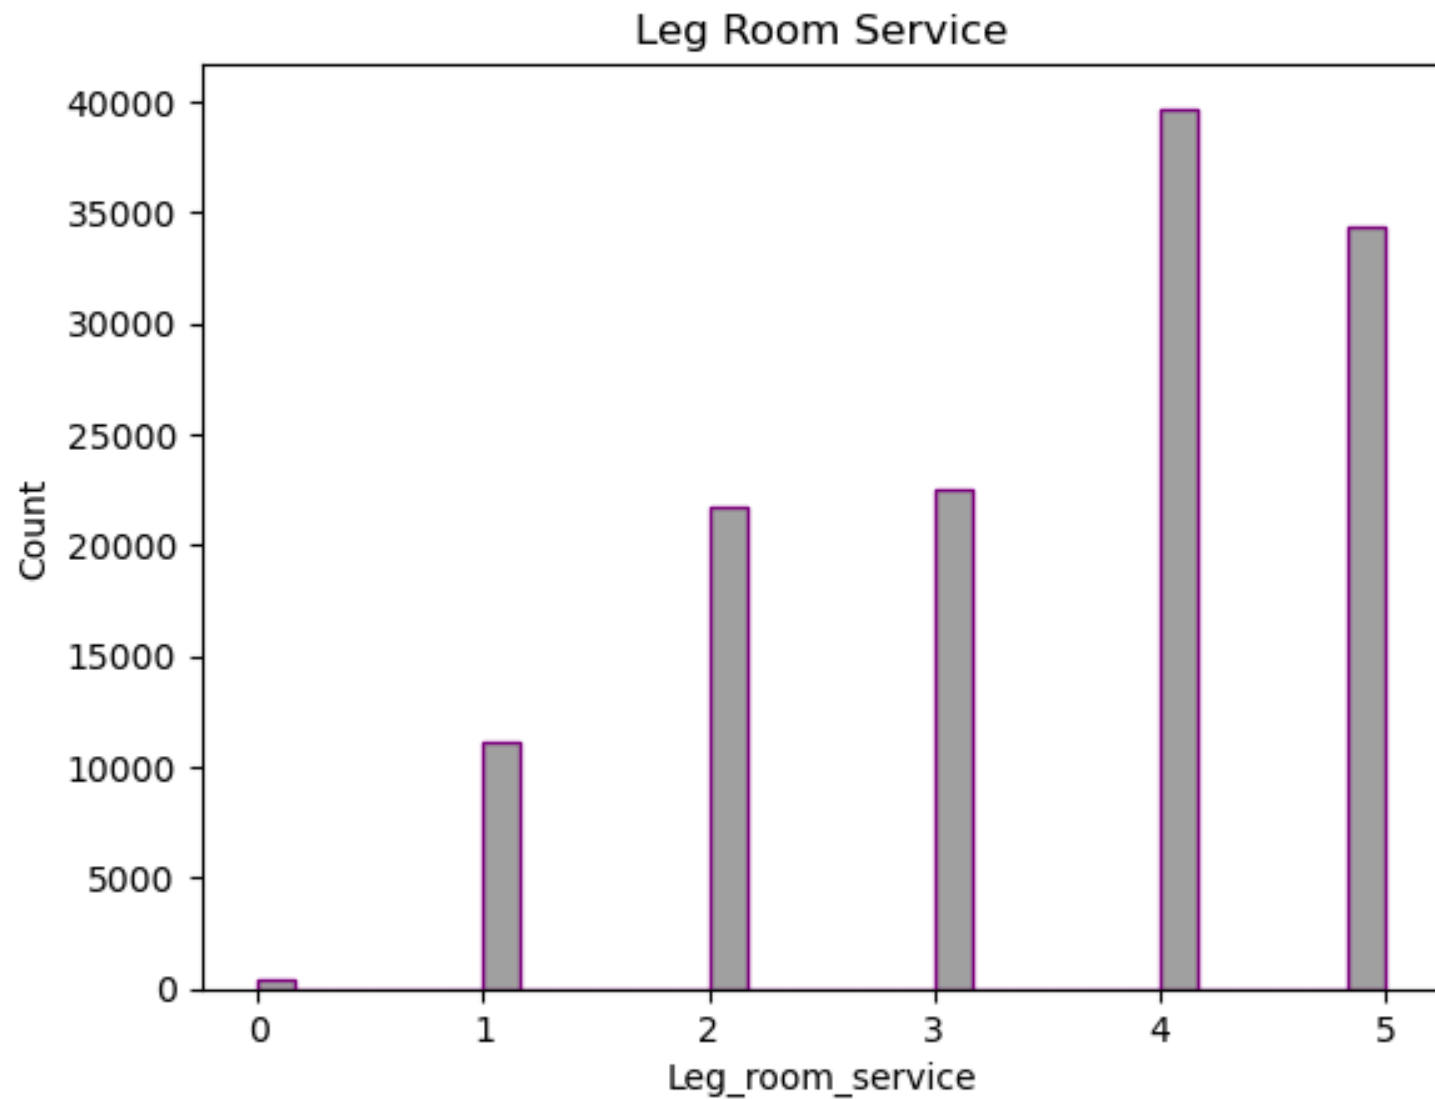
```
Out[217]: Text(0.5, 1.0, 'Food and Drink')
```



Food and Drink is a categorical variable. Food and Drink is a rating scale from 0 to 5. The rating 3 has the highest count of customers.

```
In [218... #Leg room service Histogram  
sns.histplot(num_df.Leg_room_service, bins=30, color='grey', edgecolor='purple')  
plt.title('Leg Room Service')
```

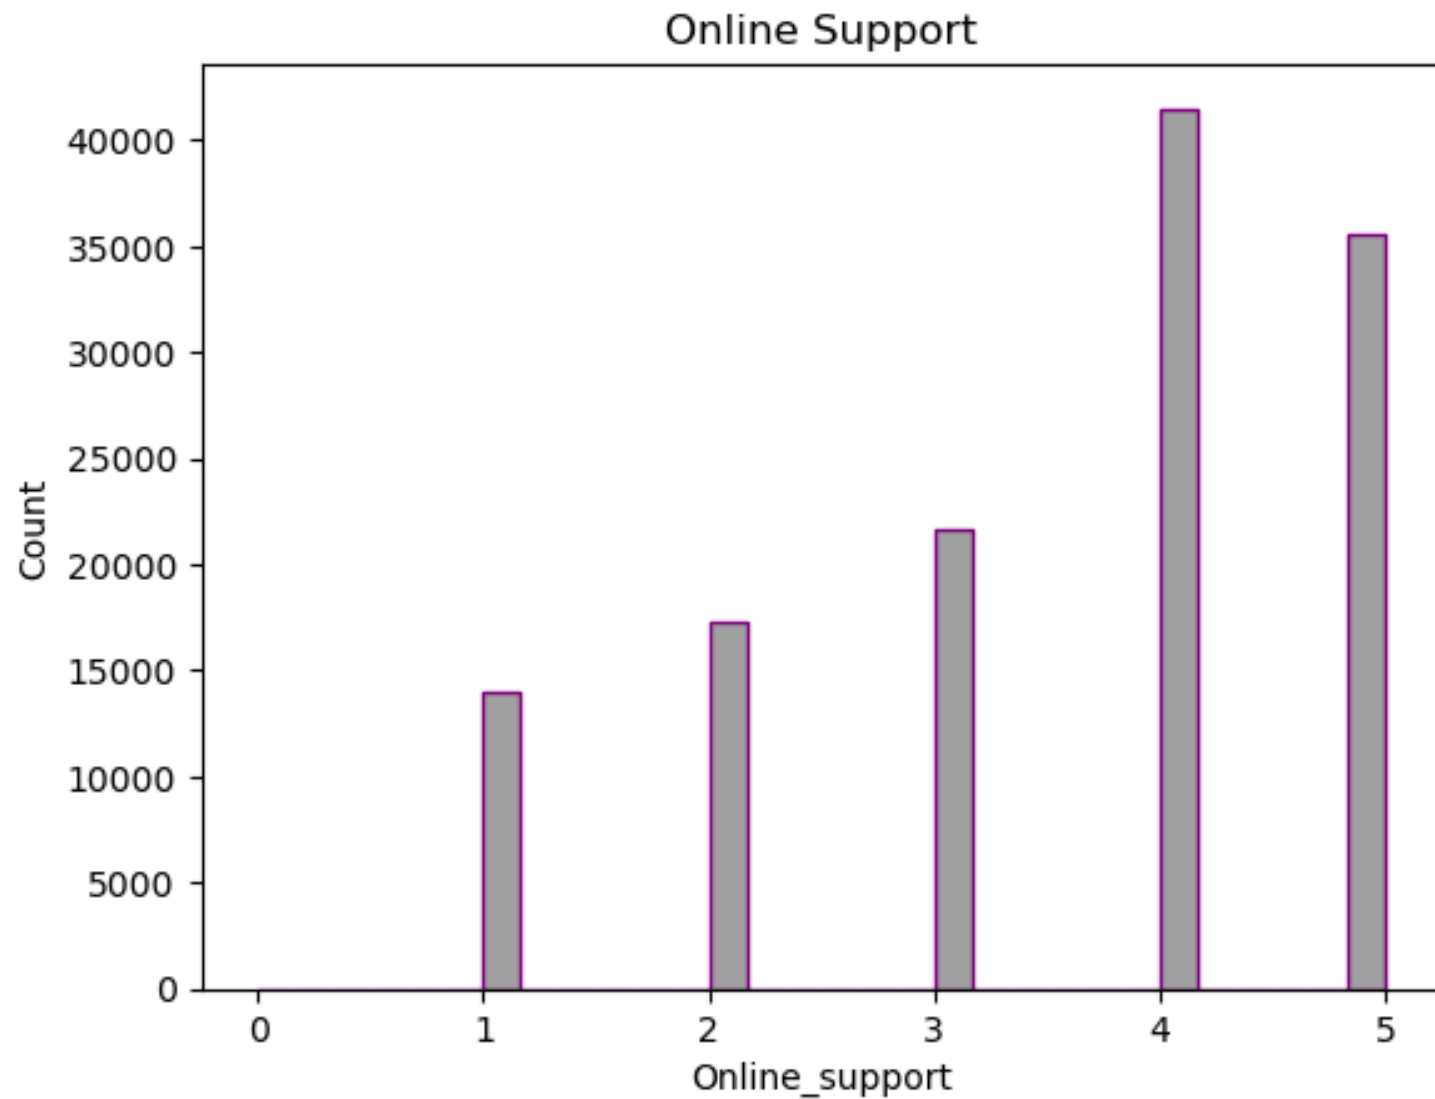
```
Out[218]: Text(0.5, 1.0, 'Leg Room Service')
```



Leg Room Service is a categorical variable. Leg Room Service is a rating scale from 0 to 5. The rating 4 has the highest count of customers.

```
In [220... #Online Support Histogram  
sns.histplot(num_df.Online_support, bins=30, color='grey', edgecolor='purple')  
plt.title('Online Support')
```

```
Out[220]: Text(0.5, 1.0, 'Online Support')
```



Online Support is a categorical variable. Online Support is a rating scale ranging from 0 to 5. There are no customer ratings for Online Support. The rating 4 is the highest customer rating.

Categorical Data

In [107...

```
#Create dataframe with only categorical variables
cat_df = df2.select_dtypes(exclude = np.number)

cat_df.columns = cat_df.columns.str.replace(' ', '_')
```

```
#Review categories  
cat_df.head()
```

```
Out[107]:
```

	satisfaction	Customer_Type	Type_of_Travel	Class
0	satisfied	Loyal Customer	Personal Travel	Eco
1	satisfied	Loyal Customer	Personal Travel	Business
2	satisfied	Loyal Customer	Personal Travel	Eco
3	satisfied	Loyal Customer	Personal Travel	Eco
4	satisfied	Loyal Customer	Personal Travel	Eco

```
In [108... #Check Data types  
cat_df.dtypes
```

```
Out[108]: satisfaction      object  
Customer_Type      object  
Type_of_Travel      object  
Class               object  
dtype: object
```


```
In [109... #Change Object to Category  
cat_df.satisfaction = cat_df.satisfaction.astype('category')  
cat_df.Customer_Type = cat_df.Customer_Type.astype('category')  
cat_df.Type_of_Travel = cat_df.Type_of_Travel.astype('category')  
cat_df.Class = cat_df.Class.astype('category')  
  
#Check Data types again  
cat_df.dtypes
```

```
Out[109]: satisfaction      category  
Customer_Type      category  
Type_of_Travel      category  
Class               category  
dtype: object
```




```
In [110]: sns.countplot(cat_df, x="satisfaction")
```

```
Out[110]: <Axes: xlabel='satisfaction', ylabel='count'>
```



```
In [111]: sns.countplot(cat_df, x="Class")
```

```
Out[111]: <Axes: xlabel='Class', ylabel='count'>
```



Logistic Regression

In [114...

```
df3.columns
```

```
Out[114]: Index(['satisfaction', 'Customer Type', 'Age', 'Type of Travel', 'Class',  
      'Flight Distance', 'Seat comfort', 'Departure/Arrival time convenient',  
      'Food and drink', 'Gate location', 'Inflight wifi service',  
      'Inflight entertainment', 'Online support', 'Ease of Online booking',  
      'On-board service', 'Leg room service', 'Baggage handling',  
      'Checkin service', 'Cleanliness', 'Online boarding',  
      'Departure Delay in Minutes', 'Arrival_Delay_Minutes'],  
      dtype='object')
```

```
In [201... #split data into features and target  
X = df3[['Cleanliness', 'Seat comfort', 'Flight Distance', 'Age', 'Departure/Arrival time convenient',  
y = df3['satisfaction']
```

```
In [202... #Split data into training and testing sets  
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.2, random_state = 42)
```

```
In [203... #Standardize  
scaler = StandardScaler()  
X_train = scaler.fit_transform(X_train)  
X_test = scaler.transform(X_test)
```

```
In [204... # Initialize and train Linear Regression model  
lr_model = LogisticRegression()  
lr_model.fit(X_train, y_train)
```

```
Out[204]: ▾ LogisticRegression  
LogisticRegression()
```

```
In [205... #Predict on test set  
y_pred = lr_model.predict(X_test)
```


```
In [206... #Evaluate Model (Confusion Matrix)  
cnf_matrix = metrics.confusion_matrix(y_test, y_pred)  
cnf_matrix
```

```
Out[206]: array([[ 8795,  2842],
        [ 2334, 11403]], dtype=int64)
```

```
In [207... #Visualize predicted and actual values
class_names=[0,1] # name of classes
fig, ax = plt.subplots()
tick_marks = np.arange(len(class_names))
plt.xticks(tick_marks, class_names)
plt.yticks(tick_marks, class_names)
# create heatmap
sns.heatmap(pd.DataFrame(cnf_matrix), annot=True, cmap="YlGnBu" ,fmt='g')
ax.xaxis.set_label_position("top")
plt.tight_layout()
plt.title('Confusion matrix', y=1.1)
plt.ylabel('Actual label')
plt.xlabel('Predicted label')
```

```
Out[207]: Text(0.5, 427.9555555555555, 'Predicted label')
```

Confusion matrix



In [208...

```
#Confusion Matrix Metrics  
target_names = ['Dissatisfied', 'Satisfied']  
print(classification_report(y_test, y_pred, target_names=target_names))
```

	precision	recall	f1-score	support
Dissatisfied	0.79	0.76	0.77	11637
Satisfied	0.80	0.83	0.82	13737
accuracy			0.80	25374
macro avg	0.80	0.79	0.79	25374
weighted avg	0.80	0.80	0.80	25374

In []: